

Business Analytics

Cabbage Price Prediction

Team 4

18102087 Lee Yeongju

20102122 Jeong Hyoan

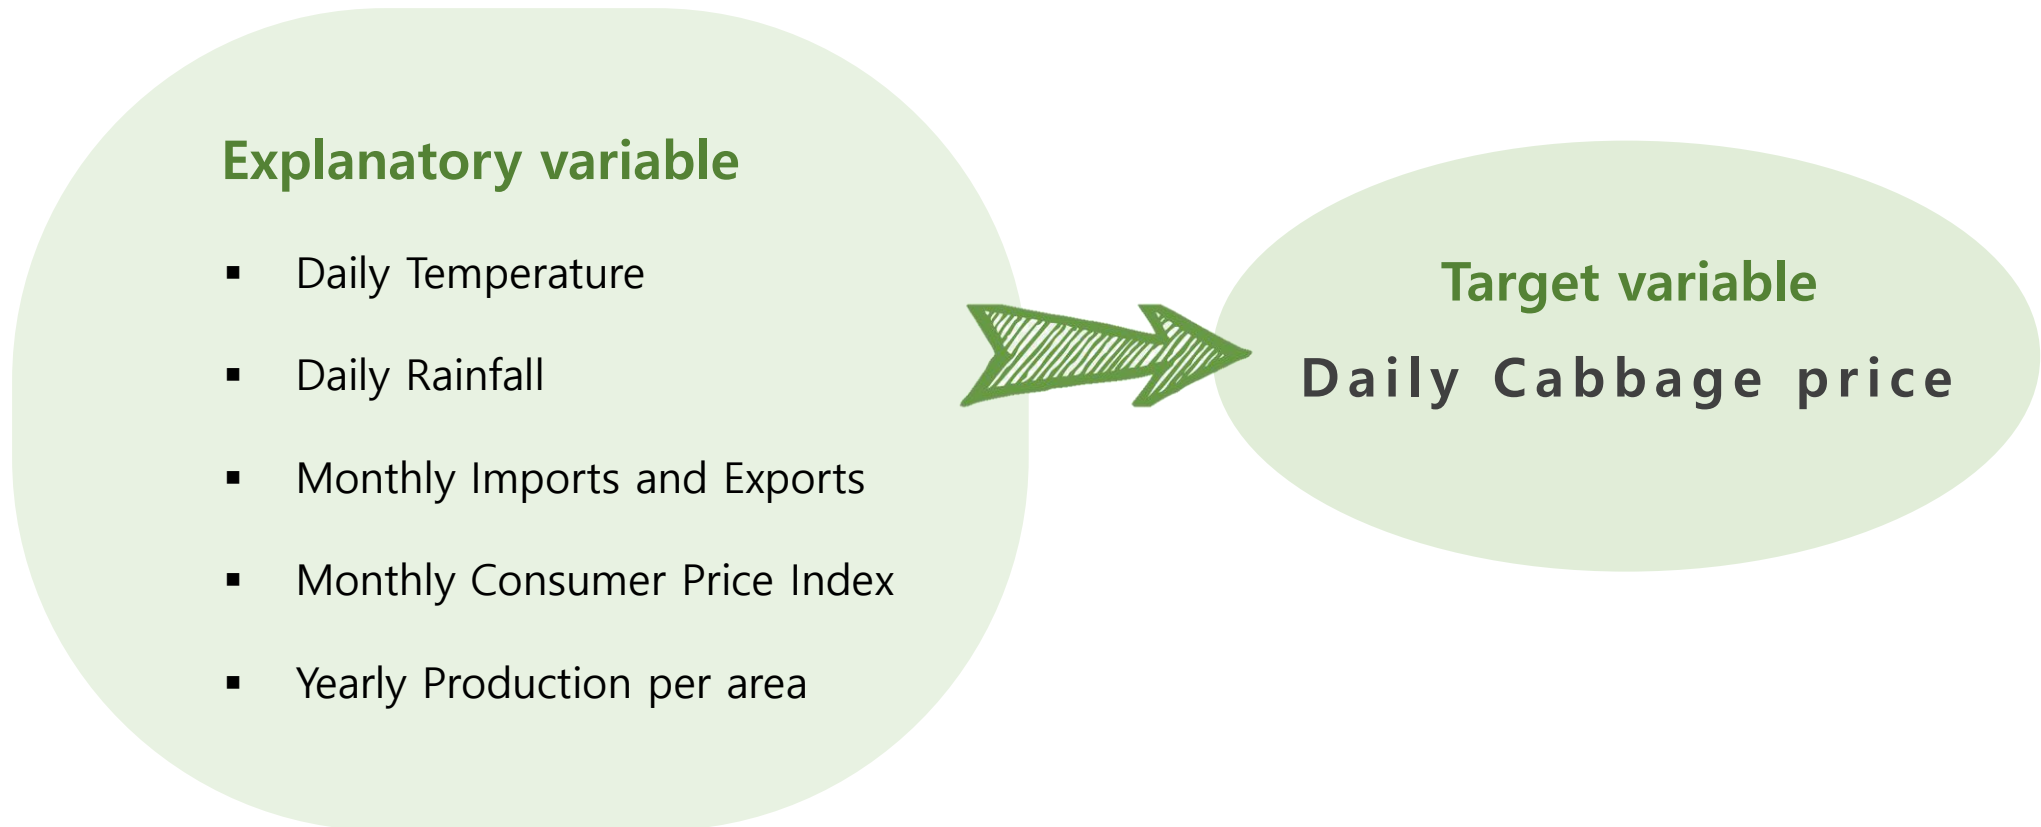
17102039 Kang Juho

Table of Contents

1. Brief explanation about topic
2. Additional preprocessing
3. Model selection
4. Feature engineering - poly
5. Hyperparameter tuning
6. Model evaluation & Comparison
7. Feature importance
8. Conclusion

1. Brief Explanation about Topic

- **Topic:** Predicting the price of cabbage, to help consumers make decisions
- **Data**



2. Additional Preprocessing – Deleting error data

- Monthly Imports and Exports data (m_trade)

→ We found error data

✓ Import amount > 0, but the imports = 0, because of unit '\$1000'

```
In [12]: m_trade
```

```
Out[12]:
```

	년월	수출액(\$1000)	수출량(kg)	수입액(\$1000)	수입량(kg)
0	2022.09	33	12522	627	1081765
1	2022.08	34	14287	436	781725
2	2022.07	296	341857	33	74030
3	2022.06	968	1362508	0	0
4	2022.05	1889	2276613	0	0
...
124	2012.05	861	1317991	38	111135
125	2012.04	54	29408	524	865030
126	2012.03	99	71207	48	135850
127	2012.02	102	179261	0	0
128	2012.01	342	644146	0	78

```
#수입액 결측치 제거
```

```
m_trade.rename(columns={'수입액($1000)': '수입액'}, inplace = True)
```

```
m_trade.rename(columns={'수입량(kg)': '수입량'}, inplace = True)
```

```
m_trade['수입액'] = m_trade['수입액'].astype(int)
```

```
m_trade['수입량'] = m_trade['수입량'].astype(int)
```

```
m_trade = m_trade.loc[(m_trade.수입액 != 0) | (m_trade.수입량 == 0)]
```

2. Additional Preprocessing – Missing values

- Checking missing values

```
In [37]: d2.isnull().sum()
Out[37]: Year          0
        Month         0
        Day           0
        평균배추가격      0
        지점           0
        평균기온(℃)      0
        최저기온(℃)      0
        최고기온(℃)      0
        강수량(mm)       0
        배추:면적 (ha)     0
        생산량 (톤)       212
        소비자물가지수      9
        수출액($1000)    189
        수출량(kg)       189
        수입액($1000)    189
        수입량(kg)       189
dtype: int64
```

- **D2 data:** 2012.01~2022.11
- **Production data:** 2012 ~ 2011
- **Consumer Price Index data:** 2012.01 ~ 2022.10
- **Imports and exports data:** 2012.01~2022.09
→ Also, it has other missing values in the data

2. Additional Preprocessing – Deleting Missing value

- Monthly Consumer Price Index (m_price)

- The dataset has missing value in the data

```
# 수출수입 데이터 결측치 제거  
d2.dropna(subset=['수출액($1000)'], inplace=True)  
d2.isnull().sum()
```

- We delete the missing data

→ Deleting null data after 2022.09 : Do not need to handle 'Consumer Price Index data'

	Year	Month	Day	평균배추 가격	지 점	평균기 온(°C)	최저기 온(°C)	최고기 온(°C)	강수량 (mm)	배추·면적 (ha)	생산량 (톤)	소비자물가 지수	수출액 (\$1000)	수출량 (kg)	수입액 (\$1000)	수입량 (kg)
20	2012	2	1	3560	전북	-6.5	-10.6	-0.4	0.3	29524	2001642.0	91.588	102.0	179261.0	0.0	0.0
21	2012	2	2	3740	전북	-10.2	-13.9	-5.7	0.1	29524	2001642.0	91.588	102.0	179261.0	0.0	0.0
22	2012	2	3	3840	전북	-7.0	-13.7	-0.9	0.0	29524	2001642.0	91.588	102.0	179261.0	0.0	0.0
23	2012	2	6	4100	전북	1.2	-4.3	6.2	1.6	29524	2001642.0	91.588	102.0	179261.0	0.0	0.0
24	2012	2	7	4100	전북	-3.7	-8.5	3.0	0.0	29524	2001642.0	91.588	102.0	179261.0	0.0	0.0
...
2642	2022	9	26	27760	전북	17.8	13.9	23.4	0.0	23853	NaN	108.930	33.0	12522.0	627.0	1081765.0
2643	2022	9	27	26720	전북	18.7	13.2	26.3	0.0	23853	NaN	108.930	33.0	12522.0	627.0	1081765.0
2644	2022	9	28	25860	전북	18.4	13.5	24.8	0.0	23853	NaN	108.930	33.0	12522.0	627.0	1081765.0
2645	2022	9	29	24640	전북	18.6	13.5	26.4	0.0	23853	NaN	108.930	33.0	12522.0	627.0	1081765.0
2646	2022	9	30	23740	전북	18.9	12.7	27.6	0.0	23853	NaN	108.930	33.0	12522.0	627.0	1081765.0

2. Additional Preprocessing – Changing Missing value


- Yearly production per area data (y_produce)

→ We change missing value in y_produce to minimum value of production

```
In [19]: y_produce
```

```
Out[19]:
```

	시점	배추:면적 (ha)	생산량 (톤)
0	2022	23853	NaN
1	2021	30085	2017507.0
2	2020	30949	2242640.0
3	2019	25837	1859705.0
4	2018	31143	2391946.0
5	2017	32416	2395686.0
6	2016	24902	1793391.0
7	2015	27174	2134976.0
8	2014	32027	2538804.0
9	2013	32020	2299251.0
10	2012	29524	2001642.0



#생산량 결측치 처리: min값으로 대체

```
production_min = d2['생산량 (톤)'].min()  
d2.fillna(production_min, inplace=True)  
d2.isnull().sum()
```

3. Model Selection

- **Linear Regression**
 - Multi-linear regression models tend to be overfitted
 - Generalization error is increased about new data
 - By solving this problem, we used Ridge and Lasso
 - We will find the best hyperparameters for each model
 - And we will compare the test scores of each models applying the hyperparameters

4. Feature Engineering

- Changing the "Month" column to categorical '계절' (season) column

→ Applying **OneHotEncoder** to the '계절'

```
d2['Month'] = d2['Month'].astype(str)
d2['Month'] = d2['Month'].str.replace('10', '가을')
d2['Month'] = d2['Month'].str.replace('12', '겨울')
d2['Month'] = d2['Month'].str.replace('1', '겨울')
d2['Month'] = d2['Month'].str.replace('2', '겨울')
d2['Month'] = d2['Month'].str.replace('3', '봄')
d2['Month'] = d2['Month'].str.replace('4', '봄')
d2['Month'] = d2['Month'].str.replace('5', '봄')
d2['Month'] = d2['Month'].str.replace('6', '여름')
d2['Month'] = d2['Month'].str.replace('7', '여름')
d2['Month'] = d2['Month'].str.replace('8', '여름')
d2['Month'] = d2['Month'].str.replace('9', '가을')
d2['Month'] = d2['Month'].str.replace('11', '가을')
```

→ Code changing months to seasons

- To put all numeric features on same scale, we used '**StandardScaler**'
- In order to enrich a feature representation, we added '**Polynomial Features**'

5. Hyperparameter Tunning

- **Lasso**
 - **Find Alpha & value of degree in Polynomial**

```
kfold = KFold(n_splits=5, shuffle=True, random_state=0)
scaler = StandardScaler()
```

Using KFold for cross-validation with data scaling and polynomial

```
d_settings=[]
alpha_settings=[]
avg_r2=[]
avg_mse=[]
```

```
best_r2=0
```

We will select the hyperparameter having the highest R2

```
for alpha in np.logspace(-1, 1, 50): # candidates for alpha
```

```
    print(alpha)
```

We set the 50 candidates of alpha in range from 0.1 to 10

```
    r2_val = []
    mse_val = []
    maxdegree=5
```

```
    for d in range(1, maxdegree):
```

We set the candidates of degree in range from 1 to 4

```
        print(d)
```

```
        alpha_settings.append(alpha)
        d_settings.append(d)
```

```
        poly = PolynomialFeatures(degree=d, include_bias=False)
```

5. Hyperparameter Tunning

- Lasso
 - Find Alpha & value of degree in Polynomial

```
for train_idx, val_idx in kfold.split(x_trainval, y_trainval):
```

```
    x_train = x_trainval.iloc[train_idx]  
    y_train = y_trainval.iloc[train_idx]  
    x_val = x_trainval.iloc[val_idx]  
    y_val = y_trainval.iloc[val_idx]
```

Dividing category variables and numerical variables for training and validation set

```
    x_train_cat = x_train[['계절']]  
    x_train_num = x_train[['평균기온(℃)', '최저기온(℃)', '최고기온(℃)', '강수량(mm)', '생산면적 (ha)', '생산량 (톤)', '소  
  
    x_val_cat = x_val[['계절']]  
    x_val_num = x_val[['평균기온(℃)', '최저기온(℃)', '최고기온(℃)', '강수량(mm)', '생산면적 (ha)', '생산량 (톤)', '소비자
```

```
    # For x_train_cat and x_val_cat, apply onehotencoding
```

```
    ohe = OneHotEncoder(sparse=False)
```

```
    ohe.fit(x_train_cat)
```

Transforming categorical features into numeric using OneHotEncoding

```
    x_train_cat_ohe = ohe.transform(x_train_cat)
```

```
    x_val_cat_ohe = ohe.transform(x_val_cat)
```

```
    # For x_train_num and x_val_num, apply standardscaler
```

```
    scaler = StandardScaler()
```

```
    scaler.fit(x_train_num)
```

```
    x_train_num_scaled = scaler.transform(x_train_num)
```

```
    x_val_num_scaled = scaler.transform(x_val_num)
```

```
    poly.fit(x_train_num_scaled)
```

```
    x_train_poly = poly.transform(x_train_num_scaled)
```

```
    x_val_poly = poly.transform(x_val_num_scaled)
```

5. Hyperparameter Tunning

- **Lasso**
 - **Find Alpha & value of degree in Polynomial**

```
# concatenate them again into x_train and X_val. # Concatenating into x_train_trans and x_val_trans  
x_train_trans = np.concatenate([x_train_cat_ohc, x_train_poly], axis=1)  
x_val_trans = np.concatenate([x_val_cat_ohc, x_val_poly], axis=1)
```

```
# training is performed with the Lasso set to the current alpha.  
lasso = Lasso(alpha = alpha, random_state=0, max_iter=10000)  
lasso.fit(x_train_trans, y_train)
```

```
# get y_valid_hat with the trained model & store r2 score in scores_val  
y_val_hat = lasso.predict(x_val_trans)
```

```
# store r2 score, mse  
r2_val.append(r2_score(y_val, y_val_hat))
```

Best score on validation set: 0.7552897
Best hyperparameters: {'alpha': 10.0, 'degree': 4}

```
mean_r2 = np.mean(r2_val) # get the cross-validation score
```

```
# When the mean_score is higher than current best score, best_score is updated and the hyperparameter at that time is saved  
if mean_r2 > best_r2:  
    best_r2 = mean_r2  
    best_parameters = {'alpha': alpha, 'degree': d}
```

```
print("Best score on validation set: {:.7f}".format(best_r2))  
print("Best hyperparameters: {}".format(best_parameters))
```

5. Hyperparameter Tunning

- Ridge
 - Find Alpha & value of degree in Polynomial

```
x_train_num_scaled = scaler.transform(x_train_num)
x_val_num_scaled = scaler.transform(x_val_num)

poly.fit(x_train_num_scaled)
x_train_poly = poly.transform(x_train_num_scaled)
x_val_poly = poly.transform(x_val_num_scaled)

# concatenate them again into x_train and X_val.
x_train_trans = np.concatenate([x_train_cat_ohe, x_train_poly], axis=1)
x_val_trans = np.concatenate([x_val_cat_ohe, x_val_poly], axis=1)

# training is performed with the Lasso set to the current alpha.
ridge = Ridge(alpha = alpha, random_state=0, max_iter=10000)
ridge.fit(x_train_trans, y_train)

# get y_valid_hat with the trained model & store r2 score in scores_val
y_val_hat = ridge.predict(x_val_trans)

# store r2 score, mse
r2_val.append(r2_score(y_val, y_val_hat))
```

```
mean_r2 = np.mean(r2_val) # get the cross-validation score
```

```
# When the mean_score is higher than current best score, best_score is updated and the hyperparameter at that time is saved
if mean_r2 > best_r2:
    best_r2 = mean_r2
    best_parameters = {'alpha': alpha, 'degree': d}
```

```
print("Best score on validation set: {:.7f}".format(best_r2))
print("Best hyperparameters: {}".format(best_parameters))
```

Best score on validation set: 0.7328941

Best hyperparameters: {'alpha': 7.543120063354615, 'degree': 3}

6. Model evaluation & Comparison

- Lasso

Best hyperparameters: {'alpha': 10.0, 'degree': 4}

```
x_trainval_cat = x_trainval[['계절']]
x_trainval_num = x_trainval[['평균기온(℃)', '최저기온(℃)', '최고기온(℃)', '강수량(mm)', '생산면적 (ha)', '생산량 (톤)', '소비자물가지수']]

x_test_cat = x_test[['계절']]
x_test_num = x_test[['평균기온(℃)', '최저기온(℃)', '최고기온(℃)', '강수량(mm)', '생산면적 (ha)', '생산량 (톤)', '소비자물가지수']]

ohe = OneHotEncoder(sparse=False)
ohe.fit(x_trainval_cat)

x_trainval_cat_ohe = ohe.transform(x_trainval_cat)
x_test_cat_ohe = ohe.transform(x_test_cat)

scaler = StandardScaler()

scaler.fit(x_trainval_num)

x_trainval_num_scaled = scaler.transform(x_trainval_num)
x_test_num_scaled = scaler.transform(x_test_num)

poly = PolynomialFeatures(degree=4, include_bias=False)

poly.fit(x_trainval_num_scaled)
x_trainval_poly = poly.transform(x_trainval_num_scaled)
x_test_poly = poly.transform(x_test_num_scaled)
```

```
In [24]: lasso = Lasso(alpha=10, random_state=0, max_iter=10000)
lasso.fit(x_trainval_trans, y_trainval)

y_test_hat = lasso.predict(x_test_trans)

test_score = r2_score(y_test, y_test_hat)

print("Test-set score:", test_score)
```

Test-set score: 0.8709727352447184

6. Model evaluation & Comparison

- **Ridge**

Best hyperparameters: {'alpha': 7.543120063354615, 'degree': 3}

```
scaler = StandardScaler()
scaler.fit(x_trainval_num)

x_trainval_num_scaled = scaler.transform(x_trainval_num)
x_test_num_scaled = scaler.transform(x_test_num)

poly = PolynomialFeatures(degree=3, include_bias=False)

poly.fit(x_trainval_num_scaled)
x_trainval_poly = poly.transform(x_trainval_num_scaled)
x_test_poly = poly.transform(x_test_num_scaled)
```

```
In [29]: ridge = Ridge(alpha = 7.543120063354615, random_state=0, max_iter=10000)
ridge.fit(x_trainval_trans, y_trainval)

y_test_hat = ridge.predict(x_test_trans)

test_score = r2_score(y_test, y_test_hat)

print("Test-set score:", test_score)

Test-set score: 0.8784319654985169
```

- ✓ **Ridge has higher test score than Lasso**
- ✓ **We are able to get a Ridge model with quite an explanatory power**

THANK YOU